

Hagelloch 1861 measles: GLCT vs Boxcar DDSA comparison (M_θ analogue)

Sandra Montes (@slmontes)

2026-03-25

Introduction

We apply discrete dynamic survival analysis (DDSA) to the 1861 measles outbreak in Hagelloch, Germany (188 cases over 86 days). The line list of this outbreak was used on an individual-level analysis by Neal & Roberts (2004, hereafter NR), which we use as a reference point.

This study treats the Hagelloch measles outbreak as a stress test of our stratified mean-field DDSA model. The Hagelloch data are a challenging case because transmission is strongly shaped by individual-level contact structure (household and classroom links), which is explicitly represented in the Neal and Roberts framework. By contrast, our specification intentionally aggregates transmission into three interacting groups and therefore does not attempt to reproduce the full contact network.

Our goal is not to reproduce NR's individual-level household/contact likelihood, but to assess what a mean-field DDSA analogue can recover from the case line list with a simplified model: the qualitative hierarchy of transmission parameters (specifically $\beta_{C1} > \beta_{C2}$ and $\beta_H > 0$, together with correct epidemic sequencing of the three groups) and the susceptible group sizes, which NR treat as fixed but DDSA infers via final-size constraints. Because the contact structure is aggregated into mean-field terms, we expect per-group attack rates to be underestimated relative to the data.

The dataset is available in R's outbreaks package as `measles_hagelloch_1861`. Individual prodrome-onset dates and rash-onset dates are recorded for all cases, allowing us to:

1. Construct the infection-time histogram (`th`) from prodrome-onset dates.
2. Compute observed infectious-period durations (`delta`) as $\delta_i = (t_{ERU,i} - t_{PRO,i}) + 4$ days, following Neal & Roberts (2004) with $d_0 = 3$ (child remains infectious 3 days after rash onset) and $d_1 = 1$ (child becomes infectious 1 day before prodrome onset), so the $+4 = d_0 + d_1$.

The infectious period has empirical mean ≈ 8 days. We model it with a NegBin discrete phase-type (DPH) distribution (rather than a simple geometric period).

Neal & Roberts (2004) partitions the 188 cases into three groups based on school enrolment:

- *Class 0 (not yet enrolled)*: No initial cases; infected via household transmission, peaked later (infection times 8–28 days).
- *Class 1 (younger school class)*: No initial cases; concentrated wave peaking around days 21–28.
- *Class 2 (older school class)*: Contains index case; infection times span the full epidemic window (days 0–40), with the main wave peaking around day 28–30.

We use the four transmission parameters from NR’s model M_θ : within-class 1 (β_{C1}), within-class 2 (β_{C2}), a household-route term from school to pre-school (β_H), and a global uniform-mixing term (β_G). The infectious period is shared across classes and modelled with a NegBin DPH. Inference uses a Gibbs scheme: β_{C1} , β_{C2} , β_H , β_G , and the calendar mean $E[\delta]$ are updated jointly by one NUTS step; while the phase count r and population sizes N_0 , N_1 , N_2 are updated by Metropolis–Hastings. We set $p = r/(E[\delta] - \delta_{\min} + r)$ given $E[\delta]$ and r . Calendar infectious period δ includes NR’s fixed +4 days, so we prepend $(\delta_{\min} - r)$ deterministic day-states to that DPH (with $\delta_{\min} = \min_i \delta_i = 4$ here), giving support $\{\delta_{\min}, \delta_{\min} + 1, \dots\}$ and $E[\delta] = (\delta_{\min} - r) + r/p$.

NR parameterise the per-day rate at which infectious individual i infects susceptible j as

$$\alpha_{ij} = \beta_H \mathbf{1}_{\{\text{same household}\}} + \beta_C^1 \mathbf{1}_{\{L_i=L_j=1\}} + \beta_C^2 \mathbf{1}_{\{L_i=L_j=2\}} + \beta_G,$$

where the global term β_G replaces the distance-decay $\beta_G \exp(-\theta \rho(i, j))$ of the full model M by setting $\theta = 0$. Our mean-field analogue translates this individual-level rate into group-level forces of infection:

$$\lambda_0 = (1 - e^{-[\beta_H(I_1+I_2)+\beta_G(I_0+I_1+I_2)]})S_0$$

$$\lambda_1 = (1 - e^{-[\beta_{C1} I_1 + \beta_G(I_0+I_1+I_2)]})S_1$$

$$\lambda_2 = (1 - e^{-[\beta_{C2} I_2 + \beta_G(I_0+I_1+I_2)]})S_2$$

Here S_k and I_k are the susceptible and infectious fractions of group k (each ranging over $[0, 1]$), and N_0 , N_1 , N_2 are inferred via Binomial final-size likelihoods.

We treat Class 0 as downstream of the school epidemic via the household-route term β_H and allow mixing between the school classes through β_G . No spatial coordinates enter the model ($\theta = 0$ in the NR notation).

We fit the model using two function-map implementations (GLCT and Boxcar) that share the same likelihood terms and recovery model. The agreement between them is used as an implementation check.

Libraries

```
using Random, LinearAlgebra
import Statistics: mean, std, median, quantile
using Distributions
using Turing
# Compatibility patch for Turing Gibbs/MH warmup copy behavior (needed for some versions)
let Ti = Turing.Inference
    if isdefined(Ti, :StoreLinkedValues)
        Ty = getfield(Ti, :StoreLinkedValues)
        @eval Base.copy(::$Ty) = $Ty()
    end
    if isdefined(Ti, :StoreUnspecifiedPriors)
        Ty = getfield(Ti, :StoreUnspecifiedPriors)
        @eval Base.copy(x::$Ty) = $Ty(copy(x.vns_with_proposal))
    end
end
using StatsPlots, Plots
using ForwardDiff
using CSV, DataFrames
using MCMCDiagnosticTools

include("./PhaseTypeDistributions.jl")
import .PhaseTypeDistributions
```

```
Random.seed!(1234)
```

Data

```
df = CSV.read("hagelloch.csv", DataFrame)

# CSV.jl return type IDs / class labels as Int
function _line_list_id(v)
    ismissing(v) && return nothing
    if v isa Union{Integer, AbstractFloat}
        x = Float64(v)
        isfinite(x) || return nothing
        return Int(round(x))
    end
end
```

```

    s = strip(string(v))
    (isempty(s) || uppercase(s) in ("NA", "N/A", ".")) && return nothing
    return tryparse{Int, s}
end

df0 = filter(r -> coalesce(_line_list_id(r.class), -1) == 0, df)
df1 = filter(r -> coalesce(_line_list_id(r.class), -1) == 1, df)
df2 = filter(r -> coalesce(_line_list_id(r.class), -1) == 2, df)

K0 = nrow(df0)
K1 = nrow(df1)
K2 = nrow(df2)

# Census-based bounds from Neal & Roberts (2004): infer N0,N1,N2 from K0,K1,K2 with N_total
N_total = 197
N0_max = N_total - K1 - K2    # = 197 - 30 - 68 = 99
N1_max = N_total - K0 - K2    # = 197 - 90 - 68 = 39
N2_max = N_total - K0 - K1    # = 197 - 90 - 30 = 77

# Shared epidemic window across all three classes
Tmax = maximum(vcat(df0.t_infection, df1.t_infection, df2.t_infection))

# Shared recovery intervals (all three groups)
delta = vcat(df0.inf_period, df1.inf_period, df2.inf_period)

# Buffer beyond Tmax: a couple of mean infectious periods
nsteps = Tmax + 2 * round{Int, mean(delta)}

# Infection-time histograms (1-indexed)
th0 = zeros{Int, nsteps}
th1 = zeros{Int, nsteps}
th2 = zeros{Int, nsteps}
for row in eachrow(df0); idx = row.t_infection + 1; idx <= nsteps && (th0[idx] += 1); end
for row in eachrow(df1); idx = row.t_infection + 1; idx <= nsteps && (th1[idx] += 1); end
for row in eachrow(df2); idx = row.t_infection + 1; idx <= nsteps && (th2[idx] += 1); end

# Support condition for offset NegBin DPH: r >= r_min.
nb_r_max = minimum(delta)

```


K0=90 (class 0), K1=30 (class 1), K2=68 (class 2)
 nsteps = 102 (covers epidemic window to day 86 + buffer)
 Recovery intervals: mean=7.95 d, min=4, max=17, CV=0.219

nb_r_max = 4 (geometric CV would be 1.0; empirical CV=0.22 → NegBin DPH warranted)

```
# Data-derived: source of infection for class 0 (pre-school) cases
class_lookup = Dict{Int, Int}()
for row in eachrow(df)
    kid = _line_list_id(row.case_ID)
    kid === nothing && continue
    cid = _line_list_id(row.class)
    cid === nothing && continue
    class_lookup[kid] = cid
end
c0_from_school = count(eachrow(df0)) do row
    inf = row.infector
    iid = _line_list_id(inf)
    iid === nothing && return false
    cls = get(class_lookup, iid, nothing)
    cls !== nothing && cls in (1, 2)
end
c0_from_c0 = count(eachrow(df0)) do row
    inf = row.infector
    iid = _line_list_id(inf)
    iid === nothing && return false
    cls = get(class_lookup, iid, nothing)
    cls !== nothing && cls == 0
end
c0_unresolved = K0 - c0_from_school - c0_from_c0
```

Class 0 infection source (from `infector` column):
from school siblings (class 1 or 2): 53 / 90 (58.9%)
from class 0 (community play): 34 / 90 (37.8%)
unresolved (missing or unknown infector): 3 / 90 (3.3%)

Plot of the infection-time histograms extracted from the data:



Models

Both the GLCT and Boxcar maps implement the same three-group M_θ force-of-infection structure and the same NegBin DPH infectious-period distribution. They differ only in how the infectious-period distribution is propagated in time: the GLCT updates a phase-vector by matrix multiplication, while the Boxcar tracks infection-age cohorts using a discrete hazard vector derived from that same DPH.

GLCT

The SIR map tracks susceptible fractions S_k and phase-vector $\mathbf{x}_k$ for each group $k \in \{0, 1, 2\}$. Transmission follows the forces of infection λ_k previously defined, and there is no feedback from class 0 to school classes.

```
"""GLCT model (M_? analogue).
Propagates three phase-vectors using the exponential rate-to-probability force of infection
function run_sir_mtheta(T_mat, alpha, beta_C1, beta_C2, beta_H, beta_G,
                        rho0, rho1, rho2, nsteps)
    T_el = promote_type(eltype(T_mat), eltype(alpha),
                        typeof(beta_C1), typeof(beta_C2),
                        typeof(beta_H),  typeof(beta_G),
                        typeof(rho0),  typeof(rho1),  typeof(rho2))

    S0 = one(T_el) - T_el(rho0);  x0 = T_el.(alpha) .* T_el(rho0)
    S1 = one(T_el) - T_el(rho1);  x1 = T_el.(alpha) .* T_el(rho1)
    S2 = one(T_el) - T_el(rho2);  x2 = T_el.(alpha) .* T_el(rho2)

    S0_hist = T_el[S0]
    S1_hist = T_el[S1]
    S2_hist = T_el[S2]
```

```

for _ in 1:nsteps
    I0 = sum(x0); I1 = sum(x1); I2 = sum(x2)
    I_total = I0 + I1 + I2

    # Forces of infection (M_0)
    # p = 1 - exp(-rate·t) with t=1; clamp to [0, Sk] for numerical safety
    r0 = T_el(beta_H) * (I1 + I2) + T_el(beta_G) * I_total
    r1 = T_el(beta_C1) * I1 + T_el(beta_G) * I_total
    r2 = T_el(beta_C2) * I2 + T_el(beta_G) * I_total
    I0 = min(max((one(T_el) - exp(-r0)) * S0, zero(T_el)), S0)
    I1 = min(max((one(T_el) - exp(-r1)) * S1, zero(T_el)), S1)
    I2 = min(max((one(T_el) - exp(-r2)) * S2, zero(T_el)), S2)

    x0 = T_mat' * x0 .+ T_el.(alpha) .* I0
    x1 = T_mat' * x1 .+ T_el.(alpha) .* I1
    x2 = T_mat' * x2 .+ T_el.(alpha) .* I2
    S0 -= I0; S1 -= I1; S2 -= I2

    push!(S0_hist, S0)
    push!(S1_hist, S1)
    push!(S2_hist, S2)
end

function to_pmf(S_hist)
    f = T_el[max(S_hist[t] - S_hist[t+1], zero(T_el)) for t in 1:nsteps]
    sF = sum(f)
    sF > T_el(1e-12) ? f ./ sF : fill(one(T_el)/nsteps, nsteps)
end

tau0 = S0_hist[1] - S0_hist[end]
tau1 = S1_hist[1] - S1_hist[end]
tau2 = S2_hist[1] - S2_hist[end]

return (tau0=tau0, f0=to_pmf(S0_hist),
        tau1=tau1, f1=to_pmf(S1_hist),
        tau2=tau2, f2=to_pmf(S2_hist))
end

```

Boxcar

The Boxcar model uses infection-age cohort bookkeeping: newly infected individuals enter age class 1 and advance one step per day, exiting at the discrete hazard $h(k) = P(W = k \mid W \geq k)$ derived from the same offset NegBin DPH. Both representations share the same marginal recovery-time distribution and the same $1 - e^{-\text{rate}}$ discretisation of the force of infection.

The hazard vector $\mathbf{h}$ is computed by iterating the transient-state vector $\mathbf{x}_{k+1} = S^\top \mathbf{x}_k$ (starting from $\mathbf{x}_1 = \alpha$) and setting $h(k) = (\mathbf{x}_k \cdot \mathbf{s}) / \sum \mathbf{x}_k$, where $\mathbf{s}$ is the absorption probability vector. The implementation is ForwardDiff-compatible (Dual-valued parameters), other AD methods may also work but have not been tested here.

```

"""Compute discrete hazard h[k] = P(W=k|W≥k) from a DPH by iterating the state vector.
If dph has ForwardDiff Dual parameters, returns a Dual-typed vector."""
function hazard_vector_from_dph(dph, K_age::Integer)
    T = eltype(dph.β)
    s = dph.s
    h = Vector{T}(undef, K_age)
    x = copy(dph.β) # state probability vector at infection-age 1
    for k in 1:K_age
        sf_val = sum(x)
        pmf_val = dot(x, s)
        h[k] = (sf_val > T(1e-15) && isfinite(sf_val)) ? pmf_val / sf_val : T(0.01)
        x = dph.S' * x # advance to next infection age
    end
    h[end] = one(T) # boundary: everyone at the last class recovers
    return clamp.(h, zero(T), one(T))
end

"""Three-group SIR Boxcar map (M_β analogue).
Propagates three infection-age vectors using discrete hazard h_rec and the same exponent
function run_sir_mtheta_boxcar(h_rec, K_age::Int,
                               beta_C1, beta_C2, beta_H, beta_G,
                               rho0, rho1, rho2, nsteps::Int)
    T_el = promote_type(eltype(h_rec), typeof(beta_C1), typeof(beta_C2),
                        typeof(beta_H), typeof(beta_G),
                        typeof(rho0), typeof(rho1), typeof(rho2))

    S0 = one(T_el) - T_el(rho0)
    S1 = one(T_el) - T_el(rho1)
    S2 = one(T_el) - T_el(rho2)

    # All initial infecteds placed in age class 1 (consistent with offset DPH β[1]=1)

```

```

x0 = zeros(T_el, K_age); x0[1] = T_el(rho0)
x1 = zeros(T_el, K_age); x1[1] = T_el(rho1)
x2 = zeros(T_el, K_age); x2[1] = T_el(rho2)

S0_hist = T_el[S0]; S1_hist = T_el[S1]; S2_hist = T_el[S2]

for _ in 1:nsteps
    I0 = sum(x0); I1 = sum(x1); I2 = sum(x2)
    I_total = I0 + I1 + I2

    # Exponential rate-to-probability discretisation – identical to run_sir_mtheta
    r0 = T_el(beta_H) * (I1 + I2) + T_el(beta_G) * I_total
    r1 = T_el(beta_C1) * I1 + T_el(beta_G) * I_total
    r2 = T_el(beta_C2) * I2 + T_el(beta_G) * I_total
    r0 = min(max((one(T_el) - exp(-r0)) * S0, zero(T_el)), S0)
    r1 = min(max((one(T_el) - exp(-r1)) * S1, zero(T_el)), S1)
    r2 = min(max((one(T_el) - exp(-r2)) * S2, zero(T_el)), S2)

    x0_new = zeros(T_el, K_age)
    x1_new = zeros(T_el, K_age)
    x2_new = zeros(T_el, K_age)

    # New infections enter age class 1
    x0_new[1] += r0; x1_new[1] += r1; x2_new[1] += r2

    # Advance age classes; h[k] fraction absorbs (recovers), 1-h[k] advances
    @inbounds for k in 1:(K_age - 1)
        h_k = T_el(h_rec[k])
        x0_new[k+1] += x0[k] * (one(T_el) - h_k)
        x1_new[k+1] += x1[k] * (one(T_el) - h_k)
        x2_new[k+1] += x2[k] * (one(T_el) - h_k)
    end
    # k == K_age: h[K_age]=1, all remaining mass is absorbed

    x0 = x0_new; x1 = x1_new; x2 = x2_new
    S0 = max(zero(T_el), S0 - r0)
    S1 = max(zero(T_el), S1 - r1)
    S2 = max(zero(T_el), S2 - r2)

    push!(S0_hist, S0); push!(S1_hist, S1); push!(S2_hist, S2)
end

```

```

function to_pmf(S_hist)
    f = T_el[max(S_hist[t] - S_hist[t+1], zero(T_el)) for t in 1:nsteps]
    sF = sum(f)
    sF > T_el(1e-12) ? f ./ sF : fill(one(T_el)/nsteps, nsteps)
end

tau0 = S0_hist[1] - S0_hist[end]
tau1 = S1_hist[1] - S1_hist[end]
tau2 = S2_hist[1] - S2_hist[end]

return (tau0=tau0, f0=to_pmf(S0_hist),
        tau1=tau1, f1=to_pmf(S1_hist),
        tau2=tau2, f2=to_pmf(S2_hist))
end

```

K_age_boxcar = 40 (99.99th-percentile truncation for Boxcar age dimension)

Inference

Inference uses a Gibbs scheme: $\beta_{C1}, \beta_{C2}, \beta_H, \beta_G, E[\delta]$ are updated jointly by one NUTS step, while the four discrete parameters (nb_r, N_0, N_1, N_2) are sampled by Metropolis–Hastings (NegBin–DPH with Neal–Roberts offset). The epidemic is seeded by a single index case in Class 2 ($\rho_2 = 1/N_2$) with no initial infections in Classes 0 or 1 ($\rho_0 = \rho_1 = 0$), consistent with the Hagelloch data where the first recorded case belongs to Class 2.

The group population sizes are not recoverable from `hagelloch.csv`, which records only the 188 infected individuals. All three susceptible population sizes are therefore inferred from the data via the Binomial final-size constraints $K_k \sim \text{Binomial}(N_k, \tau_k)$ for $k \in \{0, 1, 2\}$. By contrast, in the Hagelloch analysis of Neal & Roberts (2004), population size is treated as given (they assume $m = n = 188$) within an individual event-history likelihood. Thus, their framework does not identify group denominators from aggregate case counts alone.

Each N_k is given a $\text{DiscreteUniform}(K_k, N_{k,\max})$ prior. The lower bound K_k is the minimum possible group size (cannot have fewer susceptibles than observed cases). For the upper bound we use an external population count: Pfeilsticker’s original epidemic documentation (as cited in Neal & Roberts, 2004) records 197 children under 14 in Hagelloch at the time of the outbreak. Of these, 185 were infected and are present in the line list (plus 3 teenagers aged 14–15, giving $K = 188$ total); the 12 non-infected children each had a documented reason for non-infection (7 infants with placental immunity, 3 kept in total isolation, 1 two-year-old, 1 immigrant who had previously had measles). Neal & Roberts therefore conclude that virtually all truly susceptible children were infected, and treat the total susceptible population as $N_{\text{total}} = 197$.

Assuming all 197 children belong to one of the three groups, the maximum possible size of group k is:

$$N_{k,\max} = N_{\text{total}} - K_{k'} - K_{k''}.$$

This gives $N_0 \in [90, 99]$, $N_1 \in [30, 39]$, $N_2 \in [68, 77]$, tight ranges reflecting the near-complete attack rate ($188/197 = 95\%$).

GLCT Turing model

```
function safe_normalise(f::AbstractVector{T}, n::Int) where T
    f_pos = ifelse.(isnan.(f) .| .!isfinite.(f), zero(T), max.(f, zero(T)))
    sF = sum(f_pos)
    (sF > T(1e-15) && isfinite(sF)) ? f_pos ./ sF : fill(one(T)/n, n)
end

# Infectious period: offset NegBin DPH with minimum ⌊_floor.

@model function ddsa_hagelloch_mtheta(K0::Int, K1::Int, K2::Int,
                                       N0_max::Int, N1_max::Int, N2_max::Int,
                                       th0::Vector{Int}, th1::Vector{Int},
                                       th2::Vector{Int},
                                       delta::Vector{Int},
                                       nsteps::Int,
                                       ⌊_floor::Int)

    # Priors – population sizes inferred within census-derived bounds
    N0 ~ DiscreteUniform(K0, N0_max)
    N1 ~ DiscreteUniform(K1, N1_max)
    N2 ~ DiscreteUniform(K2, N2_max)

    # Index case in class 2; rho2 seeds one case out of N2 susceptibles
    rho0 = 0.0
    rho1 = 0.0
    rho2 = 1.0 / N2

    # Exponential priors to avoid near-zero-rate initialisation.
    beta_C1 ~ Exponential(0.2)
    beta_C2 ~ Exponential(0.1)
    beta_H ~ Exponential(0.1)
    beta_G ~ Exponential(0.1)
```

```

# Reparameterise (r, p) by mean  $\mu$  for better Gibbs/NUTS geometry.
mean_μ ~ Uniform(μ_floor + 0.5, 25.0)
nb_r ~ DiscreteUniform(1, nb_r_max)
nb_p = nb_r / (mean_μ - μ_floor + nb_r)

dph = PhaseTypeDistributions.dph_negative_binomial_min_support(nb_r, nb_p, μ_floor)

sim = run_sir_mtheta(dph.S, dph.μ, beta_C1, beta_C2, beta_H, beta_G,
                    rho0, rho1, rho2, nsteps)

tau0 = clamp(sim.tau0, 0.001, 0.999)
tau1 = clamp(sim.tau1, 0.001, 0.999)
tau2 = clamp(sim.tau2, 0.001, 0.999)

f0 = safe_normalise(sim.f0, nsteps)
f1 = safe_normalise(sim.f1, nsteps)
f2 = safe_normalise(sim.f2, nsteps)

# Binomial final-size constraints for all three groups
K0 ~ Binomial(N0, tau0)
K1 ~ Binomial(N1, tau1)
K2 ~ Binomial(N2, tau2)

# Infection-time histograms
if K0 > 0 && all(isfinite, f0) && abs(sum(f0) - 1) < 0.01
  th0 ~ Multinomial(K0, f0)
else
  Turing.@addlogprob! -Inf
end
if K1 > 0 && all(isfinite, f1) && abs(sum(f1) - 1) < 0.01
  th1 ~ Multinomial(K1, f1)
else
  Turing.@addlogprob! -Inf
end
if K2 > 0 && all(isfinite, f2) && abs(sum(f2) - 1) < 0.01
  th2 ~ Multinomial(K2, f2)
else
  Turing.@addlogprob! -Inf
end

# Shared NegBin DPH recovery likelihood
n_μ = length(delta)

```

```

    for i in 1:n_?
        delta[i] ~ dph
    end
end
end

```

Boxcar Turing model

```

@model function ddsa_hagelloch_mtheta_boxcar(K0::Int, K1::Int, K2::Int,
                                              N0_max::Int, N1_max::Int, N2_max::Int,
                                              th0::Vector{Int}, th1::Vector{Int},
                                              th2::Vector{Int},
                                              delta::Vector{Int},
                                              nsteps::Int,
                                              ?_floor::Int,
                                              K_age::Int)

    N0 ~ DiscreteUniform(K0, N0_max)
    N1 ~ DiscreteUniform(K1, N1_max)
    N2 ~ DiscreteUniform(K2, N2_max)

    rho0 = 0.0; rho1 = 0.0; rho2 = 1.0 / N2

    beta_C1 ~ Exponential(0.2)
    beta_C2 ~ Exponential(0.1)
    beta_H ~ Exponential(0.1)
    beta_G ~ Exponential(0.1)
    mean_? ~ Uniform(?_floor + 0.5, 25.0)
    nb_r ~ DiscreteUniform(1, nb_r_max)
    nb_p = nb_r / (mean_? - ?_floor + nb_r)

    dph = PhaseTypeDistributions.dph_negative_binomial_min_support(nb_r, nb_p, ?_floor)
    h_rec = hazard_vector_from_dph(dph, K_age)

    sim = run_sir_mtheta_boxcar(h_rec, K_age, beta_C1, beta_C2, beta_H, beta_G,
                                rho0, rho1, rho2, nsteps)

    tau0 = clamp(sim.tau0, 0.001, 0.999)
    tau1 = clamp(sim.tau1, 0.001, 0.999)
    tau2 = clamp(sim.tau2, 0.001, 0.999)

    f0 = safe_normalise(sim.f0, nsteps)

```

```

f1 = safe_normalise(sim.f1, nsteps)
f2 = safe_normalise(sim.f2, nsteps)

K0 ~ Binomial(N0, tau0)
K1 ~ Binomial(N1, tau1)
K2 ~ Binomial(N2, tau2)

if K0 > 0 && all(isfinite, f0) && abs(sum(f0) - 1) < 0.01
  th0 ~ Multinomial(K0, f0)
else
  Turing.@addlogprob! -Inf
end
if K1 > 0 && all(isfinite, f1) && abs(sum(f1) - 1) < 0.01
  th1 ~ Multinomial(K1, f1)
else
  Turing.@addlogprob! -Inf
end
if K2 > 0 && all(isfinite, f2) && abs(sum(f2) - 1) < 0.01
  th2 ~ Multinomial(K2, f2)
else
  Turing.@addlogprob! -Inf
end

# Recovery likelihood: same offset DPH PMF
n_? = length(delta)
for i in 1:n_?
  delta[i] ~ dph
end
end

```

Sampler

Both models use an identical sampler configuration: independence Metropolis–Hastings for nb_r , one joint Metropolis–Hastings block for (N_0, N_1, N_2) , and one joint NUTS step for the five continuous parameters $(\beta_{C1}, \beta_{C2}, \beta_H, \beta_G, \text{mean_?})$ each full Gibbs sweep. Four parallel chains are run with cold starts (initial nb_r spread over $\{1, \dots, 4\}$; moderate β 's and mean_?).

GLCT

```

n_samples = 20_000
nuts_adapt = 2_000
n_chains   = 4

```

```

gibbs_sampler = Gibbs(
  :nb_r => MH(:nb_r => DiscreteUniform(1, nb_r_max)),
  (:N0, :N1, :N2) =>
    MH(:N0 => DiscreteUniform(K0, N0_max),
      :N1 => DiscreteUniform(K1, N1_max),
      :N2 => DiscreteUniform(K2, N2_max)),
  (:beta_C1, :beta_C2, :beta_H, :beta_G, :mean_?) => NUTS(nuts_adapt, 0.8; max_depth=
)

delta_bar = mean(delta)
# Spread initial ?_C2 across chains so cold starts are less likely to lock all chains in
init_params_vec = [
  (nb_r = i,
   mean_? = delta_bar,
   N0 = K0 + 5, N1 = K1 + 5, N2 = K2 + 5,
   beta_C1 = 0.2, beta_C2 = 0.05 + 0.05 * (i - 1), beta_H = 0.1, beta_G = 0.1)
  for i in 1:n_chains
]

```

```

?_floor_ip = minimum(delta)
model_mtheta = ddsa_hagelloch_mtheta(K0, K1, K2, N0_max, N1_max, N2_max,
                                     th0, th1, th2, delta, nsteps, ?_floor_ip)

# Warm-up sample to trigger JIT
_ = sample(model_mtheta, gibbs_sampler, 1;
           progress=false, initial_params=init_params_vec[1])

chain = sample(model_mtheta, gibbs_sampler, MCMCThreads(), n_samples, n_chains;
               progress=false, initial_params=init_params_vec)

```

Chains MCMC chain (20000×12×4 Array{Float64, 3}):

```

Iterations          = 1:1:20000
Number of chains     = 4
Samples per chain    = 20000
Wall duration        = 113.92 seconds
Compute duration     = 438.39 seconds
parameters           = N0, N1, N2, beta_C1, beta_C2, beta_H, beta_G, mean_?, nb_r
internals             = logprior, loglikelihood, logjoint

```

Summary Statistics

parameters	mean	std	mcse	ess_bulk	ess_tail	rhat	?
Symbol	Float64	Float64	Float64	Float64	Float64	Float64	?
N0	98.4736	0.8422	0.0239	1086.4295	NaN	1.0035	?
N1	34.0093	2.0787	0.0682	935.1152	991.6362	1.0042	?
N2	76.6068	0.7031	0.0192	1221.6998	NaN	1.0010	?
beta_C1	0.1273	0.0362	0.0020	337.5476	264.6465	1.0187	?
beta_C2	0.0315	0.0220	0.0013	223.6748	194.5032	1.0588	?
beta_H	0.0206	0.0142	0.0008	229.6769	231.4739	1.0920	?
beta_G	0.0598	0.0083	0.0005	248.8292	275.1339	1.0519	?
mean_?	8.3618	0.2241	0.0030	5425.2417	11702.9635	1.0004	?
nb_r	3.9993	0.0302	0.0003	7470.9684	NaN	1.0003	?

1 column omitted

Quantiles

parameters	2.5%	25.0%	50.0%	75.0%	97.5%
Symbol	Float64	Float64	Float64	Float64	Float64
N0	96.0000	98.0000	99.0000	99.0000	99.0000
N1	31.0000	32.0000	34.0000	35.0000	38.0000
N2	75.0000	76.0000	77.0000	77.0000	77.0000
beta_C1	0.0506	0.1045	0.1282	0.1516	0.1956
beta_C2	0.0021	0.0147	0.0271	0.0433	0.0856
beta_H	0.0023	0.0095	0.0178	0.0291	0.0543
beta_G	0.0413	0.0549	0.0607	0.0655	0.0740
mean_?	7.9386	8.2074	8.3572	8.5112	8.8144
nb_r	4.0000	4.0000	4.0000	4.0000	4.0000

Boxcar

```
gibbs_boxcar = Gibbs(
  :nb_r      => MH(:nb_r => DiscreteUniform(1, nb_r_max)),
  (:N0, :N1, :N2) =>
    MH(:N0 => DiscreteUniform(K0, N0_max),
      :N1 => DiscreteUniform(K1, N1_max),
      :N2 => DiscreteUniform(K2, N2_max)),
  (:beta_C1, :beta_C2, :beta_H, :beta_G, :mean_?) =>
    NUTS(nuts_adapt, 0.8; max_depth=12, ?_max=1000.0)
)
```

```

init_params_bc = [
    (nb_r = i,
     mean_? = delta_bar,
     N0 = K0 + 5, N1 = K1 + 5, N2 = K2 + 5,
     beta_C1 = 0.2, beta_C2 = 0.05 + 0.05 * (i - 1), beta_H = 0.1, beta_G = 0.1)
    for i in 1:n_chains
]

```

```

model_boxcar = ddsa_hagelloch_mtheta_boxcar(K0, K1, K2, N0_max, N1_max, N2_max,
                                             th0, th1, th2, delta, nsteps,
                                             ?_floor_ip, K_age_boxcar)

_ = sample(model_boxcar, gibbs_boxcar, 1;
           progress=false, initial_params=init_params_bc[1])

chain_bc = sample(model_boxcar, gibbs_boxcar, MCMCThreads(), n_samples, n_chains;
                  progress=false, initial_params=init_params_bc)

```

Chains MCMC chain (20000×12×4 Array{Float64, 3}):

```

Iterations          = 1:1:20000
Number of chains    = 4
Samples per chain   = 20000
Wall duration       = 142.32 seconds
Compute duration    = 525.31 seconds
parameters          = N0, N1, N2, beta_C1, beta_C2, beta_H, beta_G, mean_?, nb_r
internals           = logprior, loglikelihood, logjoint

```

Summary Statistics

parameters	mean	std	mcse	ess_bulk	ess_tail	rhat ?
Symbol	Float64	Float64	Float64	Float64	Float64	Float64 ?
N0	98.3954	0.8939	0.0214	1495.5551	NaN	1.0033 ?
N1	34.0935	2.1558	0.0634	1218.8205	1604.2959	1.0054 ?
N2	76.5513	0.7646	0.0184	1478.8069	NaN	1.0014 ?
beta_C1	0.1244	0.0335	0.0017	383.2430	449.3048	1.0243 ?
beta_C2	0.0293	0.0233	0.0015	204.0762	191.6891	1.0915 ?
beta_H	0.0209	0.0128	0.0007	242.4354	289.1144	1.0600 ?
beta_G	0.0603	0.0078	0.0005	253.8788	433.8472	1.0587 ?
mean_?	8.3541	0.2248	0.0029	5844.4537	11809.4201	1.0002 ?

nb_r	3.9996	0.0255	0.0002	16027.3532	NaN	1.0002
------	--------	--------	--------	------------	-----	--------

1 column omitted

Quantiles

parameters	2.5%	25.0%	50.0%	75.0%	97.5%
Symbol	Float64	Float64	Float64	Float64	Float64
N0	96.0000	98.0000	99.0000	99.0000	99.0000
N1	31.0000	32.0000	34.0000	36.0000	39.0000
N2	74.0000	76.0000	77.0000	77.0000	77.0000
beta_C1	0.0611	0.1009	0.1235	0.1468	0.1928
beta_C2	0.0007	0.0093	0.0255	0.0437	0.0837
beta_H	0.0044	0.0113	0.0180	0.0275	0.0522
beta_G	0.0419	0.0561	0.0614	0.0658	0.0726
mean_?	7.9348	8.1991	8.3471	8.5027	8.8138
nb_r	4.0000	4.0000	4.0000	4.0000	4.0000

Results

Convergence

GLCT

ESS and $\hat{R}$:

ESS/R-hat

parameters	ess	rhat	ess_per_sec
Symbol	Float64	Float64	Float64
N0	1086.4295	1.0035	2.4782
N1	935.1152	1.0042	2.1331
N2	1221.6998	1.0010	2.7868
beta_C1	337.5476	1.0187	0.7700
beta_C2	223.6748	1.0588	0.5102
beta_H	229.6769	1.0920	0.5239
beta_G	248.8292	1.0519	0.5676
mean_?	5425.2417	1.0004	12.3755
nb_r	7470.9684	1.0003	17.0420

Boxcar

Boxcar – ESS and $\hat{R}$:

ESS/R-hat

parameters	ess	rhat	ess_per_sec
Symbol	Float64	Float64	Float64
N0	1495.5551	1.0033	2.8470
N1	1218.8205	1.0054	2.3202
N2	1478.8069	1.0014	2.8151
beta_C1	383.2430	1.0243	0.7296
beta_C2	204.0762	1.0915	0.3885
beta_H	242.4354	1.0600	0.4615
beta_G	253.8788	1.0587	0.4833
mean_?	5844.4537	1.0002	11.1257
nb_r	16027.3532	1.0002	30.5103

Posterior summaries

GLCT posterior medians:

β_{C1} median = 0.1282 (mean 0.1273; within-class 1)
 β_{C2} median = 0.0271 (mean 0.0315; within-class 2)
 β_H median = 0.0178 (mean 0.0206; household: school $\rightarrow$ class 0)
 β_G median = 0.0607 (mean 0.0598; global uniform mixing)
nb_r = 4, nb_p = 0.479 (derived), mean β = 8.36 d (joint NUTS on β 's + mean_?; β_{fl})
N0 = 99 (K0=90) prior: [90, 99]
N1 = 34 (K1=30) prior: [30, 39]
N2 = 77 (K2=68) prior: [68, 77]

Observed attack rate K_k/N_k (sample-by-sample, median [95% CI]):

K0/N0 = 0.909 [0.909, 0.938]
K1/N1 = 0.882 [0.789, 0.968]
K2/N2 = 0.883 [0.883, 0.907]

Comparison with Neal and Roberts M_{β} structure:

β_{C1}/β_{C2} ratio = 4.72 (expect > 1: younger class has higher within-class rate)
 β_H/β_G ratio = 0.29 (mean-field β_H diluted over all class-0 $\times$ school pairs; expect level)

GLCT 95% credible intervals:

```

beta_C1    : (0.0506, 0.1956)
beta_C2    : (0.0021, 0.0856)
beta_H     : (0.0023, 0.0543)
beta_G     : (0.0413, 0.074)
mean_?     : (7.9386, 8.8144)
nb_p       (derived): (0.4537, 0.5039)
nb_r       : (4, 4)
N0         : (96, 99)
N1         : (31, 38)
N2         : (75, 77)

```

Boxcar posterior medians:

```

?_C1  = 0.1235
?_C2  = 0.0255
?_H   = 0.018
?_G   = 0.0614
nb_r  = 4,  nb_p = 0.479,  mean ? = 8.35 d
N0 = 99  (K0=90)
N1 = 34  (K1=30)
N2 = 77  (K2=68)

```

Observed attack rate K_k/N_k (sample-by-sample, median [95% CI]):

```

K0/N0 = 0.909  [0.909, 0.938]
K1/N1 = 0.882  [0.769, 0.968]
K2/N2 = 0.883  [0.883, 0.919]

```

The table below displays the posterior medians for all parameters, and both GLCT and Boxcar maps should recover essentially the same parameter values.

Parameter	GLCT	Boxcar

?_C1	0.1282	0.1235
?_C2	0.0271	0.0255
?_H	0.0178	0.018
?_G	0.0607	0.0614
mean ?	8.3572	8.3471
nb_p	0.4786	0.4792
N0	99.0	99.0

N1	34.0	34.0
N2	77.0	77.0




Figure 1

R0

The basic reproduction number for the three-group model is the dominant eigenvalue of the next-generation matrix K , where entry K_{ij} gives the expected number of secondary infections in group i caused by a single infectious individual in group j , introduced into a fully susceptible population.

From the M_θ force-of-infection equations the NGM is:

$$K = \text{mean_ip} \times \begin{pmatrix} \frac{N_0}{N_0} \beta_G & \frac{N_0}{N_1} (\beta_H + \beta_G) & \frac{N_0}{N_2} (\beta_H + \beta_G) \\ \frac{N_1}{N_0} \beta_G & \frac{N_1}{N_1} (\beta_{C1} + \beta_G) & \frac{N_1}{N_2} \beta_G \\ \frac{N_2}{N_0} \beta_G & \frac{N_2}{N_1} \beta_G & \frac{N_2}{N_2} (\beta_{C2} + \beta_G) \end{pmatrix}$$

where entry K_{ij} equals $(N_i/N_j) \times \beta_{ij} \times \text{mean_ip}$, with the N_i/N_j factor arising because the model state variables are per-group fractions (one infected in group j corresponds to fraction $1/N_j$).

GLCT: median = 2.18 95% CI = (2.024, 2.389)
Boxcar: median = 2.171 95% CI = (2.021, 2.377)

Posterior predictive checks

```
function posterior_predictions_mtheta(chain, n_draws=200)
  n_total = length(vec(chain[:beta_C1]))
  idx = round(Int, range(1, n_total, length=n_draws))
  tau0_v = zeros(n_draws)
  tau1_v = zeros(n_draws)
  tau2_v = zeros(n_draws)
  f0_m = zeros(nsteps, n_draws)
  f1_m = zeros(nsteps, n_draws)
  f2_m = zeros(nsteps, n_draws)
  for (k, i) in enumerate(idx)
    bc1 = vec(chain[:beta_C1])[i]
    bc2 = vec(chain[:beta_C2])[i]
    bH = vec(chain[:beta_H])[i]
    bG = vec(chain[:beta_G])[i]
    nr = round(Int, vec(chain[:nb_r])[i])
    m2 = vec(chain[:mean_2])[i]
    np = nr / (m2 - floor_ip + nr)
    n2_i = round(Int, vec(chain[:N2])[i])
    dph_i = PhaseTypeDistributions.dph_negative_binomial_min_support(nr, np, floor_ip)
    sim = run_sir_mtheta(dph_i.S, dph_i.I, bc1, bc2, bH, bG,
                        0.0, 0.0, 1.0/n2_i, nsteps)

    tau0_v[k] = sim.tau0
    tau1_v[k] = sim.tau1
    tau2_v[k] = sim.tau2
    f0_m[:, k] .= sim.f0
    f1_m[:, k] .= sim.f1
    f2_m[:, k] .= sim.f2
  end
  return tau0_v, tau1_v, tau2_v, f0_m, f1_m, f2_m
end

tau0_post, tau1_post, tau2_post, f0_ppc, f1_ppc, f2_ppc =
  posterior_predictions_mtheta(chain)
```

```
function posterior_predictions_mtheta_boxcar(chain, n_draws=200)
  n_total = length(vec(chain[:beta_C1]))
  idx = round(Int, range(1, n_total, length=n_draws))
  tau0_v = zeros(n_draws); tau1_v = zeros(n_draws); tau2_v = zeros(n_draws)
  f0_m = zeros(nsteps, n_draws)
```




Figure 2

```

f1_m    = zeros(nsteps, n_draws)
f2_m    = zeros(nsteps, n_draws)
for (k, i) in enumerate(idx)
    bc1    = vec(chain[:beta_C1])[i];  bc2 = vec(chain[:beta_C2])[i]
    bH     = vec(chain[:beta_H])[i];   bG = vec(chain[:beta_G])[i]
    nr     = round(Int, vec(chain[:nb_r])[i])
    m2     = vec(chain[:mean_2])[i]
    np     = nr / (m2 - floor_ip + nr)
    n2_i   = round(Int, vec(chain[:N2])[i])
    dph_i  = PhaseTypeDistributions.dph_negative_binomial_min_support(nr, np, floor_ip)
    h_i    = hazard_vector_from_dph(dph_i, K_age_boxcar)
    sim    = run_sir_mtheta_boxcar(h_i, K_age_boxcar, bc1, bc2, bH, bG,
                                   0.0, 0.0, 1.0/n2_i, nsteps)
    tau0_v[k] = sim.tau0; tau1_v[k] = sim.tau1; tau2_v[k] = sim.tau2
    f0_m[:, k] .= sim.f0; f1_m[:, k] .= sim.f1; f2_m[:, k] .= sim.f2
end
return tau0_v, tau1_v, tau2_v, f0_m, f1_m, f2_m
end

tau0_bc, tau1_bc, tau2_bc, f0_ppc_bc, f1_ppc_bc, f2_ppc_bc =
    posterior_predictions_mtheta_boxcar(chain_bc)

```

Recovery interval fit

Empirical mean τ = 7.95 d, DPH posterior mean τ = 8.36 d




Figure 3




Figure 4

Discussion

We developed a mean-field analogue of Neal & Roberts (2004) with grouped compartments. Neal & Roberts (2004) explicitly set $n = m = 188$ for their analysis: the susceptible population size equals the total number of infected individuals, implying a 100% attack rate across all three groups. They justify this by noting that 185 of 197 children under 14 were infected, and that the 12 non-infected children were either infants carrying placental immunity, children kept in total isolation, a two-year-old who did not attend school, or an immigrant who had previously had measles.

In our DDSA model, all three susceptible population sizes are inferred jointly via Binomial final-size constraints, with priors derived from the number of children under 14: $N_{\text{total}} = 197$ (Neal & Roberts, 2004). This yields $N_0 \in [90, 99]$, $N_1 \in [30, 39]$, $N_2 \in [68, 77]$, ranges that are consistent with Neal & Roberts' $n=m=188$ assumption (nearly 100% attack rate) while allowing the data to determine where within those bounds the posterior concentrates.

We note that the β parameter values are not directly comparable in scale: Neal & Roberts use individual-to-individual per-day rates; our mean-field formulation uses $\beta \times I \times S$. Thus, we present a qualitative comparison of the parameter structure.

Neal & Roberts' full model included a spatial distance-decay term; in the non-spatial comparison considered here, this corresponds to setting $\theta = 0$ so the global component reduces to uniform mixing. A separate distinction concerns household transmission. In the NR individual-level model, β_H applies only to genuine household pairs (a Class 0 child is exposed via β_H only to infectious school-age siblings in the same household). In our mean-field analogue, the household term enters at the group level and is applied across all (Class 0) $\times$ (Class 1 + Class 2) pairs. This difference affects the scale of β_H : to match an aggregate household contribution, the mean-field β_H is typically smaller (spread over many more potential pairs), and transmission not attributed to β_H is absorbed by β_G .

Population treatment (all three N inferred via Binomial final-size):

$$N_0 = 99 [90, 99] \rightarrow K_0 = 90/N_0 = 90.9\%$$

$$N_1 = 34 [30, 39] \rightarrow K_1 = 30/N_1 = 88.2\%$$

$$N_2 = 77 [68, 77] \rightarrow K_2 = 68/N_2 = 88.3\%$$

(NR $n=m=188$ implies 100% attack rate; DDSA infers from data within census bounds)

Qualitative checks:

1. $\hat{\tau}_{C1} > \hat{\tau}_{C2}$: true (NR: evidence for different classroom rates)
2. $\hat{\tau}_H > 0$: true (NR: strong household effect)
3. $\hat{\tau}_G > 0$: true (NR: global mixing component)
4. $P(\hat{\tau}_{C1} > \hat{\tau}_{C2}) = 0.994$

The posterior attack rate for Class 1 ($\hat{\tau}_1 \approx K_1/N_1$) is consistent with DDSA recovering both the transmission dynamics and final size for internally homogeneous groups. The underprediction of attack rates for Classes 0 and 2 ($\hat{\tau} \approx 0.74\text{--}0.77$ vs observed $\approx 0.88\text{--}0.92$) is a consequence of the

mean-field approximation applied to groups with heterogeneous internal contact structure, and the β estimates for these groups should be interpreted accordingly. A wider choice of N_{total} (i.e. relaxing the census bound of 197 children under 14) would increase the uncertainty on N_0 and N_2 but would not qualitatively change the recovered parameter ordering or the $\beta_{C1} > \beta_{C2}$ finding.

For the infectious period, `dph_negative_binomial_min_support` enforces $r \leq \delta_{\min}$ when $\delta_{\min}$ is the calendar minimum (here, 4 days), so r cannot exceed 4 in this construction. The dwell-time parameters are in any case estimated jointly with transmission. A hard upper bound on r can make residual marginal misfit interpretable as a model-class limit rather than as precise identification of r , but the fit is still consistent with the data.

```
Class 0:  K=90,  N=99 [96, 99]  K/N (median [95% CI]) = 0.909 [0.909, 0.938]  model-
implied  $\hat{R}$  = 0.772 [0.716, 0.83]
Class 1:  K=30,  N=34 [31, 38]  K/N (median [95% CI]) = 0.882 [0.789, 0.968]  model-
implied  $\hat{R}$  = 0.884 [0.822, 0.919]
Class 2:  K=68,  N=77 [75, 77]  K/N (median [95% CI]) = 0.883 [0.883, 0.907]  model-
implied  $\hat{R}$  = 0.744 [0.668, 0.803]
```

In the observed data, Class 1 infection times cluster around days 21–28 while Class 2’s main wave spans days 26–40, so $P(\text{Class 1 peaks after Class 2})$ near zero. Another check is that $P(\text{Class 0 peaks after Class 1})$ should be close to 1, confirming that the household-driven pre-school epidemic is correctly identified as a downstream consequence of the school outbreak.

Mean peak days: Class 0=28.9, Class 1=26.9, Class 2=29.3

$P(\text{Class 1 peaks after Class 2}) = 0.02$

→ expected near 0: Class 1 should peak before Class 2’s main wave

$P(\text{Class 0 peaks after Class 1}) = 1.0$

→ expected near 1: pre-school wave is driven by household transmission from school children

In sum, both GLCT and Boxcar implementations produce consistent posteriors across all parameters, confirming that the two formulations are equivalent. Convergence is clean for β_{C1} ($\hat{R} \approx 1.01$ – 1.03 across the two implementations) but remains only marginal for β_{C2} , β_H , and β_G ($\hat{R} \approx 1.05$ – 1.09), reflecting the limited identifiability of the household and global terms under mean-field contact structure. Despite this, the model captures the primary qualitative dynamics by recovering a higher transmission rate in the younger school class compared to the older class with posterior probability $P(\beta_{C1} > \beta_{C2}) = 0.994$. It also identifies a positive household transmission route ($\beta_H > 0$), which characterises the pre-school epidemic as a secondary wave following the school outbreak, and gives consistent reproduction-number estimates between implementations (GLCT $R_0 = 2.18$, 95% CI [2.02, 2.39] vs. Boxcar $R_0 = 2.17$, 95% CI [2.02, 2.38], both posterior medians). Although the model underpredicts attack rates for certain classes due to aggregating heterogeneous contact structure into mean-field terms, these results do not alter the directional findings.